

Pyspark Analytics-PART ii

July 25, 2022

1 Part ii: Regional-analysis

```
[1]: # starting a SparkSession and creating a spark instance
import findspark
findspark.init()
findspark.find()
import pyspark
findspark.find()
```

```
[1]: 'C:\\spark-3.0.3-bin-hadoop2.7'
```

```
[2]: from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('Task2').getOrCreate()
```

1: Load the data file into a Spark DataFrame (1st DataFrame). Describe the structure of the created data frame

```
[3]: # Reading the CSV file DataSet and describing its structure
region1 = spark.read.csv('D:\\Datascience\\Region_info.csv', header=True,
    ↪ inferSchema=True)
region1.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|population|fertility| HIV|          CO2|BMI_male|
GDP|BMI_female|life|child_mortality|          Region|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      null|      null|null|      null|      null| null|      null|null|
null|      null|
| 34811059|      2.73| 0.1|3.328944661| 24.5962|12314| 129.9049|75.3|
29.5|Middle East & Nor...|
|      null|      null|null|      null|      null| null|      null|null|
null|      null|
| 19842251|      6.43| 2.0|1.474353388|22.25083| 7103| 130.1247|58.3|
192.0| Sub-Saharan Africa|
|      null|      null|null|      null|      null| null|      null|null|
null|      null|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
-----+-----+
only showing top 5 rows
```

```
[4]: region1.printSchema() #Structure of the Dataset
```

```
root
 |-- population: integer (nullable = true)
 |-- fertility: double (nullable = true)
 |-- HIV: double (nullable = true)
 |-- CO2: double (nullable = true)
 |-- BMI_male: double (nullable = true)
 |-- GDP: integer (nullable = true)
 |-- BMI_female: double (nullable = true)
 |-- life: double (nullable = true)
 |-- child_mortality: double (nullable = true)
 |-- Region: string (nullable = true)
```

2. Create a new DataFrame (2nd DataFrame) by removing the 'region' column

```
[5]: # delete single column and creating a new dataframe named Data
data = region1.drop('Region')
data.show(10)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|population|fertility| HIV|          CO2|BMI_male|
GDP|BMI_female|life|child_mortality|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|      null|      null|null|      null|      null| null|      null|null|
null|
| 34811059|      2.73| 0.1|3.328944661| 24.5962|12314| 129.9049|75.3|
29.5|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 19842251|      6.43| 2.0|1.474353388|22.25083| 7103| 130.1247|58.3|
192.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 40381860|      2.24| 0.5|4.785169983| 27.5017|14646| 118.8915|75.5|
15.4|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 2975029|      1.4| 0.1|1.804106217|25.35542| 7383| 132.8108|72.5|
20.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
```

```
| 21370348|      1.96| 0.1|18.01631327|27.56373|41312| 117.3755|81.5|
5.2|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 10 rows
```

```
[6]: # checking the schema after deletion of the rows
```

```
[7]: data.printSchema() # Checking if the Region column is removed
```

```
root
|-- population: integer (nullable = true)
|-- fertility: double (nullable = true)
|-- HIV: double (nullable = true)
|-- CO2: double (nullable = true)
|-- BMI_male: double (nullable = true)
|-- GDP: integer (nullable = true)
|-- BMI_female: double (nullable = true)
|-- life: double (nullable = true)
|-- child_mortality: double (nullable = true)
```

3. Use a graph, explore and describe the relationship between 'fertility' feature and 'life' feature in the 2nd DataFrame

```
[8]: # Staistical Summary of both Fertility and Life
data.select('fertility', 'life').describe().show()
```

```
+-----+-----+-----+
|summary|      fertility|      life|
+-----+-----+-----+
|  count|           139|           139|
|   mean| 3.005107913669065|69.60287769784175|
| stddev|1.6153544802816209|9.122189401943691|
|   min|           1.28|           45.2|
|   max|           7.59|           82.6|
+-----+-----+-----+
```

```
[9]: # Converting spark Dataframe to Python
```

```
import pandas as pd
data = data.toPandas()
```

```
[10]: # Removing missing Data
```


```
data.shape
data.dropna().head(3)
```

```
[10]:  population  fertility  HIV      CO2  BMI_male      GDP  BMI_female  life  \
1  34811059.0      2.73  0.1  3.328945  24.59620  12314.0    129.9049  75.3
3  19842251.0      6.43  2.0  1.474353  22.25083   7103.0    130.1247  58.3
5  40381860.0      2.24  0.5  4.785170  27.50170  14646.0    118.8915  75.5

      child_mortality
1              29.5
3             192.0
5             15.4
```

```
[11]: # Visualization of Data ( second dataframe)
data.plot.scatter(x='life',y='fertility')
```

```
[11]: <AxesSubplot:xlabel='life', ylabel='fertility'>
```



```
[12]: data = spark.read.option('header','true').csv('D:\\\\Datascience\\\\Region_info.
      ↪ csv', inferSchema=True)
type(data)
```

```
[12]: pyspark.sql.dataframe.DataFrame
```

```
[13]: # delete single column and creating a new dataframe named Data
data = data.drop('Region')
data.show(10)
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|population|fertility| HIV|          CO2|BMI_male|
GDP|BMI_female|life|child_mortality|
+-----+-----+-----+-----+-----+-----+-----+
-----+
|      null|      null|null|      null|      null| null|      null|null|
null|
| 34811059|      2.73| 0.1|3.328944661| 24.5962|12314| 129.9049|75.3|
29.5|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 19842251|      6.43| 2.0|1.474353388|22.25083| 7103| 130.1247|58.3|
192.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 40381860|      2.24| 0.5|4.785169983| 27.5017|14646| 118.8915|75.5|
15.4|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 2975029|      1.4| 0.1|1.804106217|25.35542| 7383| 132.8108|72.5|
20.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 21370348|      1.96| 0.1|18.01631327|27.56373|41312| 117.3755|81.5|
5.2|
+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 10 rows

```

```
[14]: data.na.drop(how='any')
      data.show(5)
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|population|fertility| HIV|          CO2|BMI_male|
GDP|BMI_female|life|child_mortality|
+-----+-----+-----+-----+-----+-----+-----+
-----+
|      null|      null|null|      null|      null| null|      null|null|
null|
| 34811059|      2.73| 0.1|3.328944661| 24.5962|12314| 129.9049|75.3|
29.5|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 19842251|      6.43| 2.0|1.474353388|22.25083| 7103| 130.1247|58.3|
192.0|

```

```
|      null|      null|null|      null|      null| null|      null|null|
null|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 5 rows
```

[]:

4. Use Spark SQL query to display the 'fertility' and 'life' columns in the 2nd DataFrame where 'fertility' is great than 1.0 and 'life' is greater than 70.

```
[15]: data.createOrReplaceTempView("Table")
spark.sql("SELECT fertility,life from Table where fertility>0 AND life>70")#
↳collect()
data.show(10)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|population|fertility| HIV|      CO2|BMI_male|
GDP|BMI_female|life|child_mortality|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|      null|      null|null|      null|      null| null|      null|null|
null|
| 34811059|      2.73| 0.1|3.328944661| 24.5962|12314| 129.9049|75.3|
29.5|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 19842251|      6.43| 2.0|1.474353388|22.25083| 7103| 130.1247|58.3|
192.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 40381860|      2.24| 0.5|4.785169983| 27.5017|14646| 118.8915|75.5|
15.4|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 2975029|      1.4| 0.1|1.804106217|25.35542| 7383| 132.8108|72.5|
20.0|
|      null|      null|null|      null|      null| null|      null|null|
null|
| 21370348|      1.96| 0.1|18.01631327|27.56373|41312| 117.3755|81.5|
5.2|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 10 rows
```

5. Build a linear regression model to predict life expectancy (the 'life' column) in the 2nd

DataFrame using the 'fertility' column as the predictor. Conduct performance evaluation for the model and make conclusions.

```
[16]: data.printSchema()
```

```
root
|-- population: integer (nullable = true)
|-- fertility: double (nullable = true)
|-- HIV: double (nullable = true)
|-- CO2: double (nullable = true)
|-- BMI_male: double (nullable = true)
|-- GDP: integer (nullable = true)
|-- BMI_female: double (nullable = true)
|-- life: double (nullable = true)
|-- child_mortality: double (nullable = true)
```

```
[17]: data.columns
```

```
[17]: ['population',
      'fertility',
      'HIV',
      'CO2',
      'BMI_male',
      'GDP',
      'BMI_female',
      'life',
      'child_mortality']
```

```
[18]: data
```

```
[18]: DataFrame[population: int, fertility: double, HIV: double, CO2: double,
BMI_male: double, GDP: int, BMI_female: double, life: double, child_mortality:
double]
```

```
[19]: # Removing missing values and dropping the non-numeric column Region.
data.na.drop().show(5)
```

```
+-----+-----+---+-----+-----+-----+-----+-----+
----+
|population|fertility|HIV|          CO2|BMI_male|
GDP|BMI_female|life|child_mortality|
+-----+-----+---+-----+-----+-----+-----+-----+
----+
|  34811059|      2.73|0.1|3.328944661| 24.5962|12314|  129.9049|75.3|
29.5|
|  19842251|      6.43|2.0|1.474353388|22.25083| 7103|  130.1247|58.3|
192.0|
|  40381860|      2.24|0.5|4.785169983| 27.5017|14646|  118.8915|75.5|
```

```

15.4|
| 2975029|      1.4|0.1|1.804106217|25.35542| 7383| 132.8108|72.5|
20.0|
| 21370348|      1.96|0.1|18.01631327|27.56373|41312| 117.3755|81.5|
5.2|
+-----+-----+---+-----+-----+-----+-----+-----+-----+
-----+
only showing top 5 rows

```

```

[20]: # counting the number of null Values.
      for col in data.columns:
          print(col.ljust(20), data.filter(data[col].isNull()).count())

```

```

population      139
fertility        139
HIV              139
CO2              139
BMI_male         139
GDP              139
BMI_female       139
life             139
child_mortality  139

```

```

[21]: data.select('fertility','life').summary('mean', '50%', 'max').show()

```

```

+-----+-----+-----+
|summary|      fertility|      life|
+-----+-----+-----+
|  mean|3.005107913669065|69.60287769784175|
|   50%|          2.41|          72.0|
|   max|          7.59|          82.6|
+-----+-----+-----+

```

```

[22]: # filling the missing Value with 50% life value
      data = data.na.fill(72.0)

```

```

[23]: # counting the number of null Values.
      for col in data.columns:
          print(col.ljust(20), data.filter(data[col].isNull()).count())

```

```

population      0
fertility        0
HIV              0
CO2              0
BMI_male         0
GDP              0
BMI_female       0

```



```
life                0
child_mortality     0
```

```
[24]: # Invoking VectorAssembler for grouping the required features
      from pyspark.ml.feature import VectorAssembler
```

```
[25]: featureassembler=VectorAssembler(inputCols=['fertility'],
                                         outputCol='feature')
```

```
[26]: # transform the element of the input and Independent feature column
      output=featureassembler.transform(data)
```

```
[27]: output.columns
```

```
[27]: ['population',
      'fertility',
      'HIV',
      'CO2',
      'BMI_male',
      'GDP',
      'BMI_female',
      'life',
      'child_mortality',
      'feature']
```

```
[28]: model_output= output.select("feature", "life")
```

```
[29]: model_output.show(5)
```

```
+-----+-----+
|feature|life|
+-----+-----+
| [72.0]|72.0|
| [2.73]|75.3|
| [72.0]|72.0|
| [6.43]|58.3|
| [72.0]|72.0|
+-----+-----+
only showing top 5 rows
```

Model Training using Linear Regression

```
[30]: from pyspark.ml.regression import LinearRegression
      #train_test_split
      #featuresCol will be the input column and labelCol will be the target column
      train_data, test_data= model_output.randomSplit([0.8, 0.2])
      lr=LinearRegression(featuresCol='feature', labelCol='life')
      lr=lr.fit(train_data)
```

Performance Evaluation

```
[31]: # Getting the coefficients
lr.coefficients
```

```
[31]: DenseVector([0.0276])
```

```
[32]: # getting intercepts
lr.intercept
```

```
[32]: 69.84713458608961
```

Prediction with linear model

```
[33]: pred=lr.evaluate(test_data)
```

```
[34]: pred.meanAbsoluteError, pred.meanSquaredError
```

```
[34]: (4.126167160095794, 45.68901889022429)
```

```
[35]: pred.predictions.show(5)
```

```
+-----+-----+-----+
|feature|life|      prediction|
+-----+-----+-----+
| [1.33]|75.3|69.88387569244462|
| [1.37]|80.0|69.88498068812447|
| [1.38]|68.2|69.88525693704443|
| [1.4]|72.5|69.88580943488435|
| [1.41]|80.4|69.88608568380432|
+-----+-----+-----+
only showing top 5 rows
```

6. Build a Lasso regression model to predict life expectancy (the 'life' column) in the 2nd DataFrame using all other columns as the predictor. Conduct performance evaluation for the model and make conclusions.

Lasso Regression model

```
[36]: logr = spark.read.option('header','true').csv('D:\\\\Datascience\\\\Region_info.
↪csv', inferSchema=True)
type(logr)
```

```
[36]: pyspark.sql.dataframe.DataFrame
```

```
[37]: logr.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|population|fertility| HIV|      CO2|BMI_male|
```



```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 10 rows
```

```
[39]: lasso_data.dtypes
```

```
[39]: [('population', 'int'),
      ('fertility', 'double'),
      ('HIV', 'double'),
      ('CO2', 'double'),
      ('BMI_male', 'double'),
      ('GDP', 'int'),
      ('BMI_female', 'double'),
      ('life', 'double'),
      ('child_mortality', 'double')]
```

```
[40]: # Checking for Missing Values
      for col in lasso_data.columns:
          print(col.ljust(20), lasso_data.filter(lasso_data[col].isNull()).count())
```

```
population          0
fertility            0
HIV                 0
CO2                 0
BMI_male            0
GDP                 0
BMI_female          0
life                0
child_mortality     0
```

Lasso prediction Model Building

```
[41]: # StringIndexer: Converts string categories to numerical categories.
      # Vector Assembler: Special to Spark API. We will find detail shortly.
      # Logistic regression based on Lasso regularization.
      from pyspark.ml.feature import StringIndexer, VectorAssembler
      from pyspark.ml.classification import LogisticRegression
      from pyspark.ml.evaluation import MulticlassClassificationEvaluator
      from pyspark.ml import Pipeline
      from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
      import numpy as np
      import pyspark.sql.functions as F
      from pyspark.ml.feature import OneHotEncoder, StringIndexer, StandardScaler
      from pyspark.ml import Pipeline
      from pyspark.ml.feature import VectorAssembler
      from pyspark.mllib.evaluation import MulticlassMetrics
      from pyspark.sql.types import FloatType
```

```
[42]: #VectorAssembler is a transformer that combines a given list of columns into a
      ↪single vector column.
      #It is useful for combining raw features and features generated by different
      ↪feature transformers
      from pyspark.ml.feature import VectorAssembler
      predictorsassembler =
      ↪VectorAssembler(inputCols=['population', 'fertility', 'HIV', 'CO2', 'BMI_male', 'GDP', 'BMI_femal
                                'child_mortality'],
                      outputCol='predictors')
```

```
[43]: # The .select() is a transformation function that is used to select the columns
      ↪from DataFrame and Dataset
      lasso_output = predictorsassembler.transform(lasso_data).
      ↪select('predictors', 'life')
```

```
[44]: lasso_output.show(10)
```

```
+-----+
|      predictors|life|
+-----+
| [3.4811059E7,2.73...|75.3|
| [1.9842251E7,6.43...|58.3|
| [4.038186E7,2.24,...|75.5|
| [2975029.0,1.4,0...|72.5|
| [2.1370348E7,1.96...|81.5|
| [8331465.0,1.41,0...|80.4|
| [8868713.0,1.99,0...|70.6|
| [348587.0,1.89,3...|72.2|
| [1.48252473E8,2.3...|68.4|
| [277315.0,1.83,1...|75.3|
+-----+
only showing top 10 rows
```

```
[45]: lasso_output.columns
```

```
[45]: ['predictors', 'life']
```

```
[46]: # The .select() is a transformation function used to select the columns from
      ↪DataFrame and Dataset
```

model training

```
[47]: # Now we split the training data into the train and test part(0.8, 0.2
      ↪respectively)
      train_lasso, test_lasso = lasso_output.randomSplit([0.8, 0.2],seed=42)
```

```
[48]: train_lasso.show(4, truncate=False) # displaying train data
```

```

+-----+-----+
|predictors                                |life|
+-----+-----+
| [277315.0,1.83,1.3,6.008278835,26.38439,16075.0,126.394,15.4] |75.3|
| [306165.0,2.91,2.4,1.36012592,27.02255,8293.0,120.9224,20.1]  |70.7|
| [321026.0,2.38,0.06,3.277725768,23.21991,12029.0,123.3223,16.0]|78.2|
| [348587.0,1.89,3.1,3.131921321,27.24594,24373.0,124.3862,14.5] |72.2|
+-----+-----+
only showing top 4 rows

```

```
[49]: test_lasso.show(4, truncate=False)# displaying test data
```

```

+-----+-----+
|predictors                                |life|
+-----+-----+
| [310033.0,2.12,0.3,6.821903051,27.20687,42294.0,118.7381,2.7] |82.0|
| [485079.0,1.63,0.3,22.16807969,27.43404,95001.0,122.3705,2.8] |80.7|
| [665414.0,5.05,0.06,0.178853064,22.06131,1440.0,132.1354,91.2] |62.6|
| [843206.0,2.74,0.1,1.277779556,26.53078,7129.0,127.4768,24.0] |65.7|
+-----+-----+
only showing top 4 rows

```

```
[50]: # Linear Model and Evaluation Matrix
evaluator = MulticlassClassificationEvaluator(labelCol='life',
↪metricName='accuracy')
```

```
[51]: lasso = LogisticRegression(labelCol='life')
```

```
[52]: # lasso Regression
evaluator
```

```
[52]: MulticlassClassificationEvaluator_fb37b604287b
```